

SWE 261P Project Part 1

Introduction

The Elasticsearch project is a remarkable endeavor, offering a distributed, RESTful search engine optimized for speed and relevance across a wide range of production-scale workloads. Its versatility extends to applications such as vector search, full-text search, log analysis, metrics, application performance monitoring (APM), security logs, and more. In this report, we delve into the Elasticsearch project, shedding light on its key functionalities, testing practices, and the crucial role played by systematic functional and partition testing. Furthermore, we explore two specific partitioning tests that have been integrated into the project, showcasing the commitment to software quality.

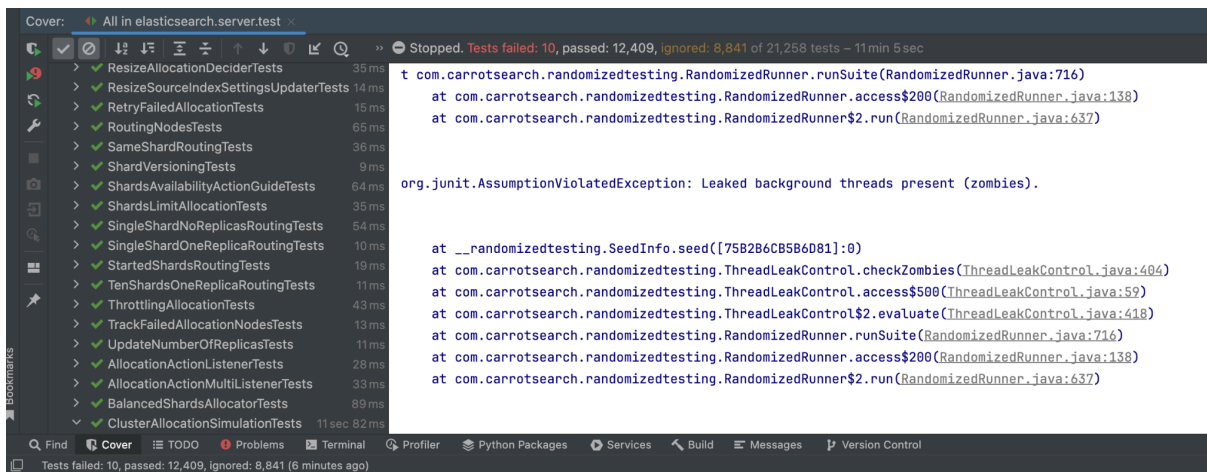
Project Overview

Elasticsearch, at its core, is designed to empower real-time search capabilities over massive datasets. While it is challenging to pinpoint the exact size of the project in terms of LOC, we can gain insights from its extensive commit history and file count. There were 592 commits with 92296 additions in the past month. The project embraces diverse technologies and languages, with Java and REST API standing out prominently. To get the project up and running, detailed instructions can be found in the provided resources, including a comprehensive guide on the build process and test execution. Alternatively, developers can easily execute tests within the popular IntelliJ IDE, making it accessible to the project's contributors.

Testing Practices

Testing forms the bedrock of software quality assurance, and Elasticsearch leaves no stone unturned. The project adopts a multi-faceted approach to testing, harnessing various frameworks tailored to different testing scenarios. These frameworks include `ESTestCase` for unit tests, `ESSingleNodeTestCase` for setting up single-node clusters, `ESIntegTestCase` for integration tests involving multiple nodes, and `ESRestTestCase` for interacting with external clusters via the REST API. Additionally, the project leverages `ESClientYamlSuiteTestCase` for YAML-based REST tests. Each framework is vital in ensuring that Elasticsearch performs flawlessly under diverse conditions.

Execution of Tests is an automated process seamlessly integrated into the development pipeline. This automation not only saves time but also enhances the reliability and consistency of the testing process, contributing significantly to the project's overall software quality.



Coverage: All in elasticsearch.server.test			
Element	Class, %	Method, %	Line, %
org.elasticsearch	72% (6234/8549)	61% (39585/63919)	61% (159355/258836)
action	81% (1054/1297)	66% (6551/9808)	67% (24765/36853)
bootstrap	70% (47/67)	54% (172/314)	40% (631/1559)
cli	92% (13/14)	67% (59/88)	70% (196/277)
client	22% (13/59)	19% (98/501)	10% (136/1266)
cluster	94% (742/787)	87% (6740/6543)	86% (26190/30384)
common	86% (746/859)	76% (4963/6491)	76% (20053/26250)
core	90% (39/43)	79% (216/273)	67% (765/1139)
discovery	94% (17/18)	88% (100/113)	93% (537/572)
env	100% (15/15)	90% (153/169)	79% (745/936)
features	100% (5/5)	100% (14/14)	100% (55/55)
gateway	98% (57/58)	85% (316/369)	84% (1953/2321)
geo	80% (4/5)	69% (52/75)	72% (240/329)
geometry	46% (23/50)	53% (215/401)	36% (612/1690)
grok	11% (2/17)	3% (3/80)	3% (10/274)
health	97% (67/69)	86% (354/411)	86% (1238/1432)
http	50% (26/52)	47% (161/337)	53% (826/1558)
index	88% (1349/1553)	79% (9970/12593)	83% (42050/50280)
indices	75% (180/239)	69% (1172/1680)	68% (5152/7554)
inference	0% (0/15)	0% (0/82)	0% (0/187)
ingest	94% (67/71)	73% (387/525)	81% (1860/2288)
internal	0% (0/1)	0% (0/1)	0% (0/1)
jdk	40% (2/5)	37% (10/27)	38% (62/159)
logging	100% (3/3)	100% (6/6)	100% (13/13)
lucene	96% (48/50)	81% (212/260)	83% (969/1156)
lz4	0% (0/8)	0% (0/55)	0% (0/313)
monitor	88% (61/69)	76% (373/488)	72% (1779/2441)
node	90% (20/22)	84% (149/176)	85% (1033/1207)
persistent	70% (33/47)	48% (146/304)	50% (450/886)
plugin	85% (6/7)	57% (8/14)	54% (13/24)
plugins	61% (44/71)	43% (169/392)	27% (398/1444)
preallocate	0% (0/10)	0% (0/21)	0% (0/81)
readiness	25% (1/4)	5% (2/39)	3% (4/130)
repositories	26% (25/96)	8% (79/925)	5% (268/4621)
reservedstate	50% (9/18)	26% (24/89)	16% (55/327)
rest	74% (183/247)	48% (589/1208)	21% (1280/5943)
script	58% (162/286)	49% (746/1500)	48% (2080/4319)
search	43% (616/1409)	28% (2868/10573)	26% (10476/39771)
secure_sm	100% (4/4)	84% (16/19)	69% (53/78)
shutdown	100% (1/1)	71% (5/7)	23% (7/30)
snapshots	36% (28/76)	22% (164/719)	16% (743/4401)
synonyms	75% (6/8)	28% (19/66)	27% (84/310)
tasks	80% (24/30)	68% (165/240)	74% (706/951)
tidytest	21% (9/42)	3% (10/286)	1% (16/1308)
telemetry	62% (31/50)	41% (79/192)	50% (147/289)
test	55% (146/262)	44% (1041/2359)	40% (3817/9454)
threadpool	85% (23/27)	74% (129/173)	73% (508/693)
transport	63% (178/280)	55% (1019/1842)	47% (3370/7043)

Systematic Testing and Partitioning

The need for systematic functional and partition testing cannot be overstated. When employed in conjunction, these testing methodologies create a robust framework for enhancing software quality.

Functional Testing, a cornerstone of quality assurance, operates as a black-box testing process, aligning test cases with the software component's specifications. In parallel, Partition Testing, or equivalence partitioning, breaks down input data into partitions from which test cases are derived. This technique efficiently covers a wide range of input conditions, ensuring comprehensive testing.

To illustrate the relevance of systematic testing in the Elasticsearch project, two specific partitioning tests have been introduced.

1. Test for Function: halfAllocatedProcessorsMaxFive

One of the critical functions within the Elasticsearch project is "halfAllocatedProcessorsMaxFive." This function plays a pivotal role as it divides the allocated processors in half while ensuring that the result is bounded between 1 and 5. This seemingly straightforward operation requires meticulous testing to ensure its accuracy and reliability.

To comprehensively test this function, partitioning techniques have been applied, creating distinct input ranges to validate its behavior:

- a. $\text{Integer.MAX_VALUE}(\text{overflow}) < \text{Input} < 1$: This partition explores scenarios where input values are less than 1. The minimal value is `Integer.MAX_VALUE` because of integer overflows when adding one to round up the results.
- b. $1 \leq \text{input} \leq 10$: Within this partition, input values ranging from 1 to 10 are tested to ensure that the function operates effectively within this defined range.
- c. $10 < \text{input} < \text{Integer.MAX_VALUE} - 1$: For input values exceeding 10 but falling short of the maximum integer value minus one, this partition examines the function's behavior under such circumstances.

2. Partitioning Test for RateLimitingFilter

Another critical component of the Elasticsearch project is the `RateLimitingFilter`, which relies on an `LruKeyCache` with a size limit of 128. When this cache reaches its predefined capacity, it automatically removes the oldest entries to maintain efficiency. To ensure the reliability of the `RateLimitingFilter`, systematic partitioning testing has been introduced.

In this specific partitioning test located at `server/src/test/java/org/elasticsearch/common/logging/RateLimitingFilterTests.java`, the partitioning strategy is defined by the cache's size limit of 128. The test creates two distinct partitions:

- a. **Size Below Limit (<128)**: This partition explores scenarios where the cache size remains below the defined limit. For instance, if we choose an input value of 50, it ensures that when accessing `key0`, this key remains in the cache, triggering a deny message (rate limit).
- b. **Size Above Limit (>128)**: In contrast, this partition tests cases where the cache size exceeds the limit. Choosing an input value of 150, it simulates a situation where `key0` would be evicted from the cache. Consequently, if we access `key0` under these conditions, it will not be found in the cache, allowing an accept message.

In addition to the partitions, boundary values have been carefully considered, further enriching the testing strategy:

- Boundary Value 127: This boundary value represents a scenario where the cache size remains below the limit. Accessing key0 in this situation will result in a denial message.
- Boundary Value 128: Here, the cache size precisely matches the defined limit. Accessing key0 will still yield a deny message, showcasing the behavior at the limit.
- Boundary Value 129: In contrast, this boundary value explores scenarios where the cache size exceeds the limit. Accessing key0 under these conditions will result in an accept message, reflecting the key's eviction from the cache.

In summary, these partitioning tests serve a vital role within the Elasticsearch project, ensuring comprehensive testing coverage across diverse input conditions. Whether it's validating the behavior of the "halfAllocatedProcessorsMaxFive" function or the efficiency of the RateLimitingFilter, systematic testing guarantees that the project maintains its high standards of software quality. As the project continues to evolve, such testing strategies remain integral in providing a reliable and efficient search engine solution.

Conclusion

In a nutshell, the Elasticsearch project stands as a testament to the commitment to software quality through systematic testing. Its versatile functionalities, diverse testing frameworks, and the incorporation of systematic functional and partition testing reflect a dedication to providing a seamless and reliable search engine solution. As the project continues to evolve, it is imperative to consider further enhancements in testing strategies, solidifying Elasticsearch's position as a leading search engine for production-scale workloads and beyond.

Reference

1. <https://github.com/jiabaow/elasticsearch/blob/main/BUILDING.md>
2. <https://github.com/jiabaow/elasticsearch/blob/main/TESTING.asciidoc>